

```

In [1]: from collections import Counter
        from scipy.sparse import csr_matrix
        from collections import defaultdict
        from sklearn.cluster import KMeans
        from numpy import *
        %matplotlib inline
        import matplotlib.pyplot as plt

        def build_matrix_1(docs, labels):
            distinctWordsAndIndex = {}
            indexIter = 0
            nnz = 0

            for idx, doc in enumerate(docs):
                frequency = doc.split()
                while frequency:
                    term, freq, *frequency = frequency
                    if term not in labels:
                        continue

                    nnz += 1
                    if term not in distinctWordsAndIndex:
                        distinctWordsAndIndex[term] = indexIter
                        indexIter += 1

            nrows = len(docs)
            ncols = len(distinctWordsAndIndex)

            ind = np.zeros(nnz, dtype=int)
            val = np.zeros(nnz, dtype=int)
            ptr = np.zeros(nrows+1, dtype=int)

            i = 0
            j = 0
            for idx, doc in enumerate(docs):
                ptr[j] = i
                j += 1
                frequency = doc.split()
                while frequency:
                    term, freq, *frequency = frequency
                    if term not in labels:
                        continue
                    ind[i] = distinctWordsAndIndex[term]
                    val[i] = int(freq)
                    i += 1
                ptr[j] = i

            mat = csr_matrix((val, ind, ptr), shape=(nrows, ncols), dtype=np.double)
            mat.sort_indices()

            return mat

```

```

In [2]: #pre-processing
        #Filter labels with length <= 3 and is present in stop words collection

        import nltk
        from nltk.corpus import stopwords

        def getValidWords():
            stop_words = stopwords.words('english')
            with open("train.clabeled", "r", encoding="utf8") as fh:
                labels = {}

```

```

    for idx, word in enumerate(fh.readlines()):
        if len(word.rstrip()) < 4 or word.rstrip() in stop_words:
            continue
        labels[str(idx+1)] = word.rstrip()
    return labels
# print(labels)

```

In [3]:

```

def BisectingKMeans(mat4, k_start, k_end, step):
    for k in range(k_start, k_end+1, step):
        print('-----')
        X = mat4
        k_list = []
        sse_list = []
        total_SSE = 0
        current_clusters = 1

        clusterMap = {}
        for idx, row in enumerate(X):
            clusterMap[idx] = idx
        finalClusterLabels = {}

        while current_clusters != k:
            kmeans = KMeans(n_clusters=2, n_init=50).fit(X)
            cluster_centers = kmeans.cluster_centers_
            sse = [0]*2
            for point, label in zip(X, kmeans.labels_):
                sse[label] += np.square(point-cluster_centers[label]).sum()
            chosen_cluster = np.argmax(sse, axis=0)
            total_SSE += sse[np.argmin(sse, axis=0)]
            chosen_cluster_data = X[kmeans.labels_ == chosen_cluster]
            newClusterMap = {}
            clusterIter = 0
            #to keep track of the index of the clusters
            for idx, x in enumerate(kmeans.labels_):
                if(x != chosen_cluster):
                    finalClusterLabels[clusterMap[idx]] = current_clusters
                elif current_clusters + 1 == k:
                    finalClusterLabels[clusterMap[idx]] = current_clusters + 1
                else:
                    newClusterMap[clusterIter] = clusterMap[idx]
                    clusterIter += 1
            clusterMap = newClusterMap
            current_clusters += 1
            assigned_cluster_data = X[kmeans.labels_ != chosen_cluster]
            X = chosen_cluster_data
            k_list.append(k)
            sse_list.append(kmeans.inertia_)
            print_internal_metrics(mat4, finalClusterLabels)
        print('-----')
    return finalClusterLabels

```

In [4]:

```

#Internal Metrics
from sklearn import metrics

def print_internal_metrics(mat, labels_dict):
    labels = [labels_dict[key] for key in sorted(labels_dict.keys())]
    n_clusters_ = len(set(labels)) - (1 if -1 in labels else 0)

    print('K: %d' % n_clusters_)
    print("Silhouette Coefficient: %0.3f" % metrics.silhouette_score(mat, labels))
    print("Calinski Harabasz Score: %0.3f" % metrics.calinski_harabasz_score(mat, labels))

```

```

In [5]: import numpy as np
from sklearn.datasets import make_blobs
from sklearn.decomposition import TruncatedSVD
from sklearn.feature_extraction.text import TfidfTransformer
from datetime import datetime

#Read the file
with open("train.dat", "r", encoding="utf8") as fh:
    rows = fh.readlines()

#Select valid words
valid_words = getValidWords()

#build csr matrix from valid words only
mat1 = build_matrix_1(rows, valid_words)

# TF-IDF transform to normalise the matrix
tfidf_transformer=TfidfTransformer(smooth_idf=True,use_idf=True)
tf_idf_vector=tfidf_transformer.fit(mat1).transform(mat1)

#print CST information
print('Shape before SVD', tf_idf_vector.shape)

print("Start Time =", datetime.now().strftime("%H:%M:%S"))

# components = 200
# while components < 501:
#     print('concepts', components)
#     wcss = []
#     for i in range(2,23,2):
#         kmeans = KMeans(n_clusters=i,init='k-means++',max_iter=300,n_init=1)
#         kmeans.fit(tf_idf_vector)
#         wcss.append(kmeans.inertia_)
#         print('k', i)
#     plt.plot(range(2,23,2),wcss)
#     plt.title('The Elbow Method')
#     plt.xlabel('Number of clusters')
#     plt.ylabel('WCSS')
#     plt.savefig('elbow.png')
#     plt.show()
#     components += 50

# This
components = 4500
while components < 4501:
    print('=====')
    print('SVD number of concepts = ', components)
    tsvd = TruncatedSVD(n_components=components)
    mat4 = tsvd.fit(tf_idf_vector).transform(tf_idf_vector)
    print('Variance ratio sum', tsvd.explained_variance_ratio_.sum())
    finalClusterLabels = BisectingKMeans(mat4, 14, 14, 2)
    components += 50
    print('=====')

print("End Time =", datetime.now().strftime("%H:%M:%S"))

```

Shape before SVD (8580, 26237)

Start Time = 17:05:21

=====

SVD number of concepts = 4500

Variance ratio sum 0.9590014631709839

-----

K: 14

Silhouette Coefficient: 0.022

Calinski Harabasz Score: 38.850

-----  
=====  
End Time = 17:23:38

```
In [6]: #write the cluster info to output file  
labels = [finalClusterLabels[key] for key in sorted(finalClusterLabels.keys())  
with open("output.dat", "w", encoding="utf8") as file:  
    for item in labels:  
        file.write("%s\n" % str(item))
```

```
In [7]: #print the number of elements in each cluster  
count_labels = {}  
for label in labels:  
    if label not in count_labels:  
        count_labels[label] = 1  
    else:  
        count_labels[label] = int(count_labels[label]) + 1  
print(count_labels)
```

```
{1: 2914, 14: 1116, 10: 160, 5: 1137, 11: 43, 12: 21, 6: 646, 4: 611, 13: 48,  
3: 654, 2: 754, 7: 301, 9: 133, 8: 42}
```